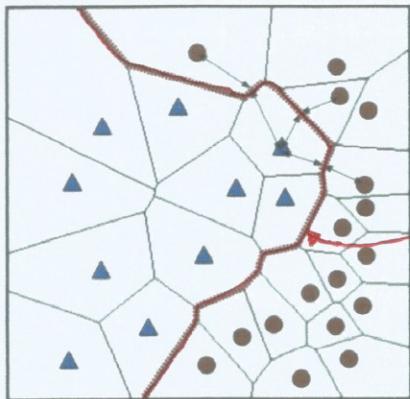


NEAREST NEIGHBOURS

VORONOI TESSELATION

→ is a partition of a plane into regions close to each of a given set of objects

These regions are called Voronoi cells



1-NN example

→ For each point there is a corresponding region consisting of all points of the plane closer to that point than to any other

DECISION BOUNDARY — Non-linear

↳ Points at same distance from 2 diff. training examples

↳ Notice that algo. is sensitive to outliers

↳ Single mislabeled example dramatically changes boundary

↳ No notion of a prior

↳ No notion tht. one class may be more frequent than another class

SOLN: — Use more than one nearest neighbour to make decision
[k-NN where $k > 1$]

k-NN CLASSIFICATION ALGO.

→ **No training phase;** the algo. gets the training set & testing set at the same time

→ Given the training examples $\{x_i, y_i\}$ and a testing point x :

1) Compute distance $D(x, x_i)$ to every training example x_i

2) Select k closest instances $x_{i1} \dots x_{ik}$ and their labels $y_{i1} \dots y_{ik}$

3) Output the class y^* which is most frequent in $y_{i1} \dots y_{ik}$

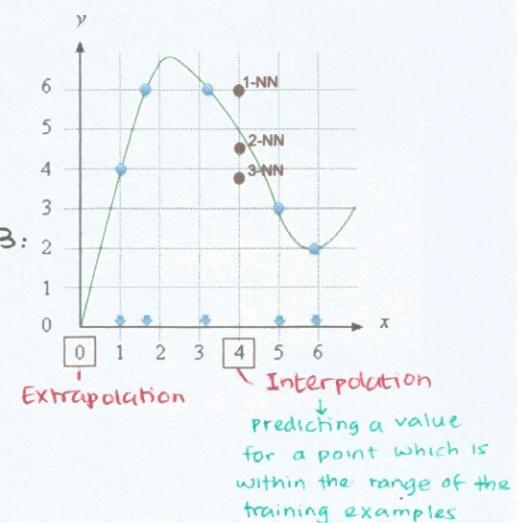
k-NN REGRESSION ALGO.

→ The y_i will now be a real-valued target

→ Algo. is similar to k-NN classification algo. but for step 3:

3) Output the mean of y_{i1}, \dots, y_{ik} : $\hat{y} = \frac{1}{K} \sum_{j=1}^K y_{ij}$

Instead of taking the most frequent class,
we'll avg. the labels



Predicting a value
for a point which is
within the range of the
training examples

SELECTING VALUE FOR K

- If value of k is too high, all the instances will be classified in the most probable class
- If too low, the DB is unstable — small changes to training set causes large changes in classification
- To select k , try different values of k and pick the k that performs best on the validation set — If we use training set, we'd always get $k=1$

PRACTICAL ISSUES

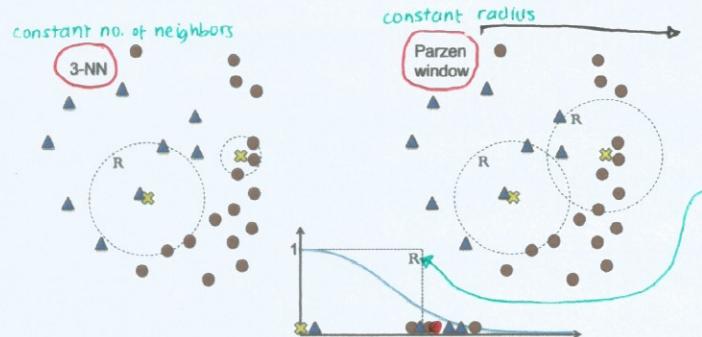
1) Resolving ties

- Ties occur when there are equal no. of votes for 2 diff. classes
- To break ties, you can
 - do it randomly (flip coin)
 - use the prior (pick class w/ greater prior)
 - use 1-NN classifier to decide

2) Missing values

- have to fill in, otherwise can't compute distance
- The value shld. affect distance as possible
- can use avg. value across entire dataset

KERNELISED NN METHOD



In the Parzen version of NN, we consider the training instances that fall in a fixed region, around the testing instance

This radius R is adding some variance to our classification.

i.e. If we make R a little bit smaller, we have to break a tie and if we make it a bit bigger, the point is now classified red.

↳ Hence we replace the 'boxy' fn w/ a fn tht. is smoother [The blue line]

KERNEL FN — Give more weights to neighbors closer to the test pt.

↳ In kernelised NN method, every training instance contributes to the answer in proportion to the kernel value for that training instance & the testing instance

↳ Removes the said variance

$$f(x) = \text{sign} \left[\sum_i y_i \cdot \underbrace{\mathbb{1}_{\|x_i - x\| \leq R}}_{\text{Kernel function}} \right] = \text{sign} \left[\sum_i y_i \cdot \underbrace{K(x_i, x)}_{\text{Kernel function}} \right]$$

Kernelised NN

→ similar form to kernelised SVM:
 $\text{sign} \left[\sum_i \alpha_i y_i K(x_i, x) \right]$
The difference lies in α_i

PROS & CONS OF k-NN

(+)

- Almost no assumptions abt. the data
 - assumptions implied by the distance fn (only locally!)
 - non-parametric approach
 - ↳ not fitting any distribution to the data

(-)

- Need to handle missing values
- Sensitive to outliers
- Sensitive to lots of irrelevant attributes
 - ↳ Like NB and unlike DTs
- Computationally expensive
 - need to store ALL training examples
 - need to compute distance to ALL examples
 - $O(nd)$ cost of computing distance
no. of training examples
 - $\uparrow n$ - system will become slower

MAKING k-NN FASTER

- Training takes $O(d)$ time but testing takes $O(nd)$, hence we can:

1) **Reduce d** through dimensionality reduction

2) **Reduce n** by not comparing to all training examples

↳ **IDEA** - Quickly identify $m \ll n$ potential nearest neighbors

↳ **K-D trees** work if you are working w/ low-dimensional, real-valued data

• $O(d \log_2 n)$, only works when $d \ll n$ and can miss neighbors

↳ **Inverted lists** work w/ high-dimensional, discrete (sparse) data

• $O(n'd')$ where $d' \ll d$, $n' \ll n$ and you don't miss neighbors

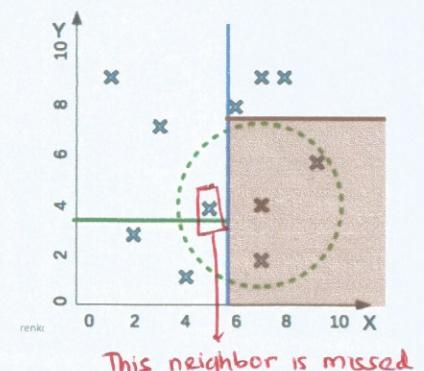
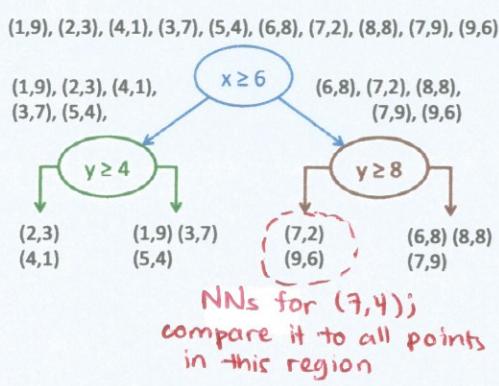
↳ **Locality-sensitive hashing** works w/ high-dimensional, real-valued/discrete data

• $O(n'd)$ where $n' \ll n$ and can miss neighbors

K-D TREES

→ Pick random dimension
 ↓
 Find median
 ↓
 Split data Repeat

e.g. Testing pt. is $(7,4)$



LOCALITY-SENSITIVE HASHING (LSH)

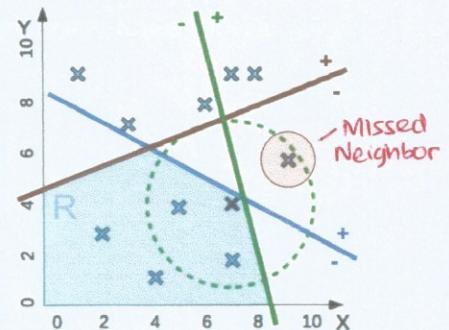
- Start by drawing random hyperplanes $h_1 \dots h_k$
 - ↳ Space will be sliced into 2^k regions (Polytopes)

- Compare testing pt. only to training pts. within the same region.

→ Complexity is $O(kd + d \frac{n}{2^k}) \approx O(d \log_2 n)$

To find the region R

Comparison to $\frac{n}{2^k}$ points in R



INVERTED LIST

- Is a data struct. used by search engines
- List all training e.g. that contain particular attribute
- ASSUMPTION: most attribute values are zero (Sparseness)
- $O(d\sqrt{n})$

D1: "send your password"	spam
D2: "send us review"	ham
D3: "send us password"	spam
D4: "send us details"	ham
D5: "send your password"	spam
D6: "review your account"	ham

new email: "account review"

↳ Given a new testing example,
merge inverted lists for attributes present

send	→	1	2	3	4	5
your	→	1	5	6		
review	→	2	6			
account	→	6				
password	→	1	3	5		

→ No chance that any NNs
will be missed